

## ▼ Spoken Digit Recognition

In this notebook, You will do Spoken Digit Recognition.

Input - speech signal, output - digit number

It contains

1. Reading the dataset. and Preprocess the data set. Detailed instructions are given below.  
You have to write the code in the same cell which contains the instruction.
2. Training the LSTM with RAW data
3. Converting to spectrogram and Training the LSTM network
4. Creating the augmented data and doing step 2 and 3 again.

Instructions:

1. Don't change any Grader Functions. Don't manipulate any Grader functions. If you manipulate any
2. Please read the instructions on the code cells and markdown cells. We will explain what to write
3. please return outputs in the same format what we asked. Eg. Don't return List if we are asking
4. Please read the external links that we are given so that you will learn the concept behind the

Saved successfully!



section if necessary, please follow them.

Every Grader function has to return True.

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

### Importing Libraries

```
import numpy as np
import pandas as pd
import librosa
import librosa.display
import os
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import seaborn as sns
import glob
from joblib import Parallel, delayed
```

```

import time
import librosa.display
from collections import Counter
from sklearn.utils import shuffle
from sklearn.metrics import confusion_matrix, f1_score, classification_report
import tensorflow as tf
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.callbacks import LearningRateScheduler, ReduceLROnPlateau, ModelChecker
from tensorflow.keras.layers import Input, BatchNormalization, Dropout, LSTM, Dense, AveragePooling2D
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.utils import plot_model
tf.keras.backend.clear_session()
##if you need any imports you can do that here.

```

## ▼ Data Preparation

```

#read the all file names in the recordings folder given by us
#(if you get entire path, it is very useful in future)
#save those files names as list in "all_files"
path = '/content/drive/MyDrive/Audio_Spoken_Digit/recordings'
all_files = []
for files in os.listdir(path):
    names = os.path.join(path, files)
    all_files.append(names)
all_files

```

Saved successfully!

```

'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_nicolas_39.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/7_theo_29.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/2_jackson_19.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/3_nicolas_38.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/5_jackson_49.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_yweweler_40.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/3_theo_26.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/5_theo_46.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/9_nicolas_38.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/2_nicolas_23.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/4_yweweler_38.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_theo_36.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_jackson_47.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/9_yweweler_35.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_yweweler_45.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/3_jackson_19.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/9_jackson_46.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_nicolas_0.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_yweweler_39.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/4_yweweler_24.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/4_nicolas_48.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/7_jackson_35.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/3_yweweler_41.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/4_yweweler_41.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/9_jackson_27.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_nicolas_6.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_theo_7.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/5_jackson_0.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/0_yweweler_15.wav'

```

```

'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_yweweler_21.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/4_theo_48.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/5_yweweler_41.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/7_yweweler_26.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/0_nicolas_48.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/5_theo_48.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_nicolas_39.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/1_theo_20.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/9_jackson_19.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_theo_13.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/7_jackson_8.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/2_jackson_42.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/2_nicolas_35.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/3_yweweler_12.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/5_nicolas_11.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/3_nicolas_3.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_jackson_6.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/0_theo_44.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_jackson_21.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/4_jackson_41.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_theo_13.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/0_jackson_14.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/7_jackson_0.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/7_jackson_38.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/0_theo_41.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/8_yweweler_30.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/1_yweweler_45.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/9_theo_31.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/6_theo_22.wav',
'/content/drive/MyDrive/Audio_Spoken_Digit/recordings/7_yweweler_24.wav'

```

## Grader function 1

Saved successfully!



```

temp = len(all_files)==2000
temp1 = all([x[-3:]=="wav" for x in all_files])
temp = temp and temp1
return temp
grader_files()

```

True

#Create a dataframe(name=df\_audio) with two columns(path, label).

#You can get the label from the first letter of name.

#Eg: 0\_jackson\_0 --> 0

#0\_jackson\_43 --> 0

label = []

for files in os.listdir(path):

label.append(files[0])

len(label)

2000

```

df_audio = pd.DataFrame(list(zip(all_files, label)), columns = ['path', 'label'])
df_audio

```

|      | path  | label |  |
|------|---|-------|--|
| 0    | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 7     |  |
| 1    | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 3     |  |
| 2    | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 0     |  |
| 3    | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 0     |  |
| 4    | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 9     |  |
| ...  | ...   | ...   |  |
| 1995 | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 2     |  |
| 1996 | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 3     |  |
| 1997 | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 3     |  |
| 1998 | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 0     |  |
| 1999 | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 4     |  |

2000 rows × 2 columns

```
#info
```

```
df_audio.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   path    2000 non-null      object
1   label   2000 non-null      object
dtypes: object(2)
memory usage: 31.4+ KB
```

Saved successfully!

```
df_audio.describe()
```

|        | path  | label |  |
|--------|---|-------|--|
| count  | 2000  | 2000  |  |
| unique | 2000  | 10    |  |
| top    | /content/drive/MyDrive/Audio_Spoken_Digit/reco... | 5     |  |
| freq   | 1   | 200   |  |

## Grader function 2

```
def grader_df():
    flag_shape = df_audio.shape==(2000,2)
    flag_columns = all(df_audio.columns==['path', 'label'])
```

```
list_values = list(df_audio.label.value_counts())
flag_label = len(list_values)==10
flag_label2 = all([i==200 for i in list_values])
final_flag = flag_shape and flag_columns and flag_label and flag_label2
return final_flag
grader_df()
```

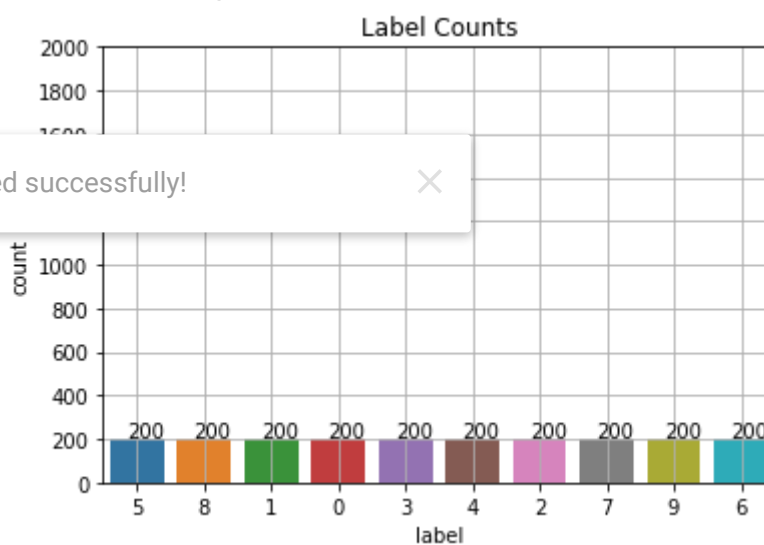
True

```
df_audio = shuffle(df_audio, random_state=33)
```

```
# Count of Labels.
total = len(df_audio)
print(Counter(df_audio['label']))
ax = sns.countplot(df_audio['label'])
for p in ax.patches:
    ax.annotate('{}' .format(p.get_height()), (p.get_x()+0.25, p.get_height()+5))
```

```
ax.yaxis.set_ticks(np.linspace(0, total, 11))
plt.grid(True)
plt.title("Label Counts")
plt.show()
```

```
Counter({'5': 200, '8': 200, '1': 200, '0': 200, '3': 200, '4': 200, '2': 200, '7': 200, '9': 200, '6': 200})
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass
FutureWarning
```



From above all labels are equally numbered.

## Train and Validation split

```
#split the data into train and validation and save in X_train, X_test, y_train, y_test
#use stratify sampling, use random state of 45, use test size of 30%
```

```
X = df_audio['path']
Y = df_audio['label'].astype("int32")
```

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.3, stratify = Y, r
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
((1400,), (600,), (1400,), (600,))
```

### Grader function 3

```
def grader_split():
    flag_len = (len(X_train)==1400) and (len(X_test)==600) and (len(y_train)==1400) and (l
    values_ytrain = list(y_train.value_counts())
    flag_ytrain = (len(values_ytrain)==10) and (all([i==140 for i in values_ytrain]))
    values_ytest = list(y_test.value_counts())
    flag_ytest = (len(values_ytest)==10) and (all([i==60 for i in values_ytest]))
    final_flag = flag_len and flag_ytrain and flag_ytest
    return final_flag
grader_split()

True
```

### Preprocessing

All files are in the "WAV" format. We will read those raw data files using the librosa

```
sample_rate = 22050
def load_wav(x, get_duration=True):
    '''This return the array values of audio with sampling rate of 22050 and Duration'''
    #loading the wav file with sampling rate of 22050
    samples, sample_rate = librosa.load(x, sr=22050)
    duration = librosa.get_duration(samples, sample_rate)
    return [samples, duration]
else:
    return samples
```

```
# Using Parallel jobs to get the samples and duration for train and test.
a = Parallel(n_jobs=-1, verbose = 1)(delayed(load_wav)(train) for train in X_train)
b = Parallel(n_jobs=-1, verbose = 1)(delayed(load_wav)(test) for test in X_test)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 174 tasks      | elapsed:    4.7s
[Parallel(n_jobs=-1)]: Done 786 tasks      | elapsed:   50.0s
[Parallel(n_jobs=-1)]: Done 1036 tasks     | elapsed:   2.1min
[Parallel(n_jobs=-1)]: Done 1386 tasks     | elapsed:   3.8min
[Parallel(n_jobs=-1)]: Done 1400 out of 1400 | elapsed:   3.9min finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done  46 tasks      | elapsed:   13.3s
[Parallel(n_jobs=-1)]: Done 196 tasks      | elapsed:   59.2s
[Parallel(n_jobs=-1)]: Done 446 tasks      | elapsed:   2.2min
[Parallel(n_jobs=-1)]: Done 600 out of 600 | elapsed:   3.0min finished
```

```
Train_samples = []
Train_duration = []
```

```
Test_samples = []
Test_duration = []
```

```
a = np.array(a)
Train_samples = a[:,0].tolist()
Train_duration = a[:,1].tolist()
a.dtype
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWar
    """Entry point for launching an IPython kernel.
dtype('O')
```

```
a = np.array(a)
Train_samples = a[:,0].tolist()
Train_duration = a[:,1].tolist()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWar
    """Entry point for launching an IPython kernel.
```

```
b = np.array(b)
Test_samples = b[:,0].tolist()
Test_duration = b[:,1].tolist()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWar
    """Entry point for launching an IPython kernel.
```

Saved successfully!

(1400, 600)

```
#use load_wav function that was written above to get every wave.
#save it in X_train_processed and X_test_processed
# X_train_processed/X_test_processed should be dataframes with two columns(raw_data, durat
```

```
X_train_processed = pd.DataFrame(list(zip(Train_samples, Train_duration)), columns = ['raw
X_test_processed = pd.DataFrame(list(zip(Test_samples, Test_duration)), columns = ['raw_da
X_train_processed.shape, X_test_processed.shape
```

((1400, 2), (600, 2))

```
del a
del b
```

```
# Train Duration.
X_train_processed['duration'].plot.hist()
```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f05fc31bf90>



# Test Duration.

```
X_test_processed['duration'].plot.hist()
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f05ff8a4e10>



Saved successfully!



with step size of 10 for train data duration.

```
for i in range(0, 101, 10):
```

```
    per = np.percentile(X_train_processed['duration'], i)
```

```
    print(i, 'th percentile is ', per)
```

```
0 th percentile is 0.1435374149659864
10 th percentile is 0.2632517006802721
20 th percentile is 0.30514285714285716
30 th percentile is 0.3346485260770975
40 th percentile is 0.35863945578231293
50 th percentile is 0.39045351473922907
60 th percentile is 0.4159546485260771
70 th percentile is 0.44731519274376413
80 th percentile is 0.481342403628118
90 th percentile is 0.5617823129251701
100 th percentile is 2.195918367346939
```

# Print 90 to 100 percentile values with step size of 1 for train data duration.

```
for i in range(90, 101):
```

```
    per = np.percentile(X_train_processed['duration'], i)
```

```
    print(i, 'th percentile is ', per)
```

```
90 th percentile is 0.5617823129251701
91 th percentile is 0.5755346938775511
```



```

92 th percentile is 0.5859174603174603
93 th percentile is 0.6012321995464855
94 th percentile is 0.6179111111111111
95 th percentile is 0.6330226757369615
96 th percentile is 0.6447981859410431
97 th percentile is 0.6635891156462584
98 th percentile is 0.6956090702947844
99 th percentile is 0.79601179138322
100 th percentile is 2.195918367346939

```

## Grader function 4

```

def grader_processed():
    flag_columns = (all(X_train_processed.columns==['raw_data', 'duration'])) and (all(X_t
    flag_shape = (X_train_processed.shape ==(1400, 2)) and (X_test_processed.shape==(600,2
    return flag_columns and flag_shape
grader_processed()

True

```

## ▼ Updating data suitably for the model.

Based on our analysis 99 percentile values are less than 0.8sec so we will limit maximum length of X\_train\_processed and X\_test\_processed to 0.8 sec. It is similar to pad\_sequence for a text dataset.

Saved successfully!



ing sampling rate of 22050 so one sec will give array of  
is  $0.8 * 22050 = 17640$

Pad with Zero if length of sequence is less than 17640 else Truncate the number.

Also create a masking vector for train and test.

Masking vector, value = 1 if it is real value, 0 if it is pad value. Masking vector data type must be bool.

```
max_length = 17640
```

```

## As discussed above, Pad with Zero if length of sequence is less than 17640 else Truncat
## Save in the X_train_pad_seq, X_test_pad_seq
## Also Create masking vector X_train_mask, X_test_mask
## All the X_train_pad_seq, X_test_pad_seq, X_train_mask, X_test_mask will be numpy arrays

```

```

X_train_mask = np.array([np.ones(g.shape[0]) for g in X_train_processed['raw_data'].values
X_train_mask.shape

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWar
"""Entry point for launching an IPython kernel.
(1400,)

```

```
X_test_mask = np.array([np.ones(g.shape[0]) for g in X_test_processed['raw_data'].values])
X_test_mask.shape
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWarning:
    """Entry point for launching an IPython kernel.
(600,)
```

```
X_train_mask[-50].shape, X_train_processed['raw_data'].values[-50].shape

((6789,), (6789,))
```

```
X_train_pad_seq = pad_sequences(X_train_processed['raw_data'], maxlen=max_length, padding=
X_train_mask = pad_sequences(X_train_mask, maxlen=max_length, padding='post', dtype = bool
```

```
X_test_pad_seq = pad_sequences(X_test_processed['raw_data'], maxlen=max_length, padding='p
X_test_mask = pad_sequences(X_test_mask, maxlen=max_length, padding='post', dtype = bool,
```

```
X_test_mask
```

```
array([[ True,  True,  True, ..., False, False, False],
       [ True,  True,  True, ..., False, False, False],
       [ True,  True,  True, ..., False, False, False],
       ...,
       [ True,  True,  True, ..., False, False, False],
       [ True,  True,  True, ..., False, False, False],
       ...,
       [ True,  True,  True, ..., False, False, False],
       [ True,  True,  True, ..., False, False, False]])
```

Saved successfully!

## Grader function 5

```
def grader_padoutput():
    flag_padshape = (X_train_pad_seq.shape==(1400, 17640)) and (X_test_pad_seq.shape==(600
    flag_maskshape = (X_train_mask.shape==(1400, 17640)) and (X_test_mask.shape==(600, 176
    flag_dtype = (X_train_mask.dtype==bool) and (X_test_mask.dtype==bool)
    return flag_padshape and flag_maskshape and flag_dtype
grader_padoutput()
```

```
True
```

```
# Changing the dimension.
```

```
X_T_pad = tf.expand_dims(X_train_pad_seq, axis = 2)
X_T_pad.shape, X_T_pad[0].shape
```

```
(TensorShape([1400, 17640, 1]), TensorShape([17640, 1]))
```

```
X_Te_pad = tf.expand_dims(X_test_pad_seq, axis = 2)
X_Te_pad.shape, X_Te_pad[0].shape
```

```
(TensorShape([600, 17640, 1]), TensorShape([17640, 1]))
```

```
Train = [X_T_pad, X_train_mask]
```

```
Val = [X_Te_pad, X_test_mask]
```

```
Train_data = (Train, y_train.values)
```

```
Val_data = (Val, y_test.values)
```

## ▼ 1. (Model 1) Giving Raw data directly.

Now we have

Train data: X\_train\_pad\_seq, X\_train\_mask and y\_train

Test data: X\_test\_pad\_seq, X\_test\_mask and y\_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes "X\_train\_pad\_seq" as input, "X\_train\_mask" as mask
2. Get the final output of the LSTM and give it to Dense layer of any size and then give
3. Use tensorboard to plot the graphs of loss and metric(use micro F1 score as metric) a

Saved successfully!



ularization

```
tf.keras.backend.clear_session()
```

```
class Metrics(tf.keras.callbacks.Callback):
```

```
    def __init__(self, x = None, y = None):
```

```
        self.training_data = x
```

```
        self.validation_data = y
```

```
    def on_train_begin(self, logs = {}):
```

```
        ## on begin of training, we are creating a instance variable called history
```

```
        self.history={'train_f1_score': [], 'val_f1_score': []}
```

```
    def on_epoch_end(self, epoch, logs = {}):
```

```
        ## on end of each epoch, we will get logs and update the self.history dict
```

```
        train_predict = self.model.predict(self.training_data[0], batch_size = 50)
```

```
        train_bin = np.argmax(train_predict, axis = 1)
```

```
        train_targ = self.training_data[1]
```

```
        _train_f1 = f1_score(train_targ, train_bin, average = 'micro')
```

```

val_predict = self.model.predict(self.validation_data[0], batch_size = 50)
val_bin = np.argmax(val_predict, axis = 1)
val_targ = self.validation_data[1]
_val_f1 = f1_score(val_targ, val_bin, average = 'micro')

self.history['val_f1_score'].append(_val_f1)
self.history['train_f1_score'].append(_train_f1)

print(' - train_f1_score : ', _train_f1, ' - val_f1_score : ', _val_f1)
return

```

```

def changeLearningRate(epoch):
    global initial_learningrate
    epoch = epoch + 1
    if epoch % 5 == 0:
        initial_learningrate *= 0.55
    return initial_learningrate

```

```

import os
save = 'model_save/*.hdf5'
r = glob.glob(save)
for i in r:
    os.remove(i)

```

```

filepath="model_save/model-{epoch:02d}-{val_sparse_categorical_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_sparse_categorical_accuracy',
                             reduce_lr_on_plateau=True, monitor = 'val_loss', factor = 0.001, patience = 1, verbose =
                             changeLearningRate, verbose=1)

```

Saved successfully!

```
metrics = Metrics(Train_data, Val_data)
```

```
callbacks = [metrics, checkpoint, reduce_lr, lrschedule]
```

```
reg = tf.keras.regularizers.L2(l2=0.01)
```

```

input_layer = Input(shape=(17640,1), name = 'input_layer')
input_mask = Input(shape=(17640,), name = 'mask_layer', dtype=bool)
ls = LSTM(units = 128, name = 'LSTM')(input_layer, mask = input_mask)
ls = BatchNormalization()(ls)
dc1 = Dense(512,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed
dc1 = Dropout(0.35)(dc1)
out = Dense(10,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal

```

```
model_raw = Model(inputs = [input_layer,input_mask], outputs = out)
```

```
model_raw.summary()
```

Model: "model"

| Layer (type)                    | Output Shape       | Param # | Connected to                          |
|---------------------------------|--------------------|---------|---------------------------------------|
| input_layer (InputLayer)        | [(None, 17640, 1)] | 0       |                                       |
| mask_layer (InputLayer)         | [(None, 17640)]    | 0       |                                       |
| LSTM (LSTM)                     | (None, 128)        | 66560   | input_layer[0][0]<br>mask_layer[0][0] |
| batch_normalization (BatchNorma | (None, 128)        | 512     | LSTM[0][0]                            |
| FC2 (Dense)                     | (None, 512)        | 66048   | batch_normalization[                  |
| dropout (Dropout)               | (None, 512)        | 0       | FC2[0][0]                             |
| FC3 (Dense)                     | (None, 10)         | 5130    | dropout[0][0]                         |
| Total params: 138,250           |                    |         |                                       |
| Trainable params: 137,994       |                    |         |                                       |
| Non-trainable params: 256       |                    |         |                                       |

```
model_raw.compile(optimizer=tf.keras.optimizers.RMSprop(0.001),
                  loss='sparse_categorical_crossentropy',
                  metrics=['sparse_categorical_accuracy'])
```

```
model_history = model_raw.fit(Train, y_train.values, batch_size = 50, validation_data = (V
```

Saved successfully!

```
====] - 22s 775ms/step - loss: 2.2931 - sparse_cat
42857143 - val_f1_score : 0.09833333333333333
```

Epoch 00005: val\_sparse\_categorical\_accuracy did not improve from 0.10500

Epoch 00005: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.  
Epoch 6/50

Epoch 00006: LearningRateScheduler reducing learning rate to 0.00055.  
28/28 [====] - 22s 771ms/step - loss: 2.2915 - sparse\_cat  
- train\_f1\_score : 0.10000000000000002 - val\_f1\_score : 0.10000000000000002

Epoch 00006: val\_sparse\_categorical\_accuracy did not improve from 0.10500

Epoch 00006: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.  
Epoch 7/50

Epoch 00007: LearningRateScheduler reducing learning rate to 0.00055.  
28/28 [====] - 21s 770ms/step - loss: 2.2876 - sparse\_cat  
- train\_f1\_score : 0.10357142857142858 - val\_f1\_score : 0.10333333333333333

Epoch 00007: val\_sparse\_categorical\_accuracy did not improve from 0.10500

Epoch 00007: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.  
Epoch 8/50

Epoch 00008: LearningRateScheduler reducing learning rate to 0.00055.  
28/28 [====] - 22s 776ms/step - loss: 2.2858 - sparse\_cat

```

28/28 [-----] - 22s 770ms/step - loss: 2.2837 - sparse_categorical_accuracy: 0.10333333333333333
- train_f1_score : 0.10285714285714286 - val_f1_score : 0.10333333333333333

Epoch 00008: val_sparse_categorical_accuracy did not improve from 0.10500

Epoch 00008: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.
Epoch 9/50

Epoch 00009: LearningRateScheduler reducing learning rate to 0.00055.
28/28 [=====] - 22s 777ms/step - loss: 2.2837 - sparse_categorical_accuracy: 0.10333333333333333
- train_f1_score : 0.10000000000000002 - val_f1_score : 0.10000000000000002

Epoch 00009: val_sparse_categorical_accuracy did not improve from 0.10500

Epoch 00009: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.
Epoch 10/50

Epoch 00010: LearningRateScheduler reducing learning rate to 0.0003025000000000000.
28/28 [=====] - 22s 783ms/step - loss: 2.2796 - sparse_categorical_accuracy: 0.10333333333333333
- train_f1_score : 0.10071428571428571 - val_f1_score : 0.10000000000000002

Epoch 00010: val_sparse_categorical_accuracy did not improve from 0.10500

Epoch 00010: ReduceLROnPlateau reducing learning rate to 3.0250000418163836e-07.
Epoch 11/50

Epoch 00011: LearningRateScheduler reducing learning rate to 0.0003025000000000000.
28/28 [=====] - 22s 776ms/step - loss: 2.2766 - sparse_categorical_accuracy: 0.10333333333333333
- train_f1_score : 0.10142857142857142 - val_f1_score : 0.10333333333333333

Epoch 00011: val_sparse_categorical_accuracy did not improve from 0.10500

```

Saved successfully!



```

['sparse_categorical_accuracy']
['train_f1_score', 'val_f1_score', 'train_sparse_categorical_accuracy']

```

```

loss = model_history.history['loss']
val_loss = model_history.history['val_loss']

f1 = metrics.history['train_f1_score']
val_f1 = metrics.history['val_f1_score']

epochs_range = range(50)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.show()




plt.plot(epochs_range, f1, label='Training micro F1 score')
plt.plot(epochs_range, val_f1, label='Validation micro F1 score')
plt.legend(loc='lower right')
plt.title('Training and Validation F1 Score')
plt.show()

```

```

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

```



```
model_raw.save('Model_raw_1_0.1517_final.h5')
```

```
model_raw = load_model('/content/model-31-0.1333.hdf5')
```

```
#write the LSTM
y_pred = model_raw.predict(Val)
y_pred = np.argmax(y_pred, axis = 1)

print('Classification Report')
print(classification_report(y_test.values, y_pred))
```

```
Classification Report
              precision    recall  f1-score   support

     0       0.07       0.03       0.05         60
     1       0.00       0.00       0.00         60
     2       0.19       0.38       0.25         60
     3       0.16       0.12       0.14         60
     4       0.14       0.02       0.03         60
     5       0.31       0.07       0.11         60
     6       0.14       0.52       0.22         60
     7       0.16       0.10       0.12         60
     8       0.21       0.23       0.22         60
     9       0.08       0.05       0.06         60


 accuracy          0.15         0.15         0.15        600
 macro avg         0.15         0.15         0.12        600
 weighted avg      0.15         0.15         0.12        600
```

```
ax= plt.subplot()
sns.heatmap(confusion_matrix(y_test.values, y_pred), annot=True, ax = ax, fmt='g', cmap='G')
```

Saved successfully!

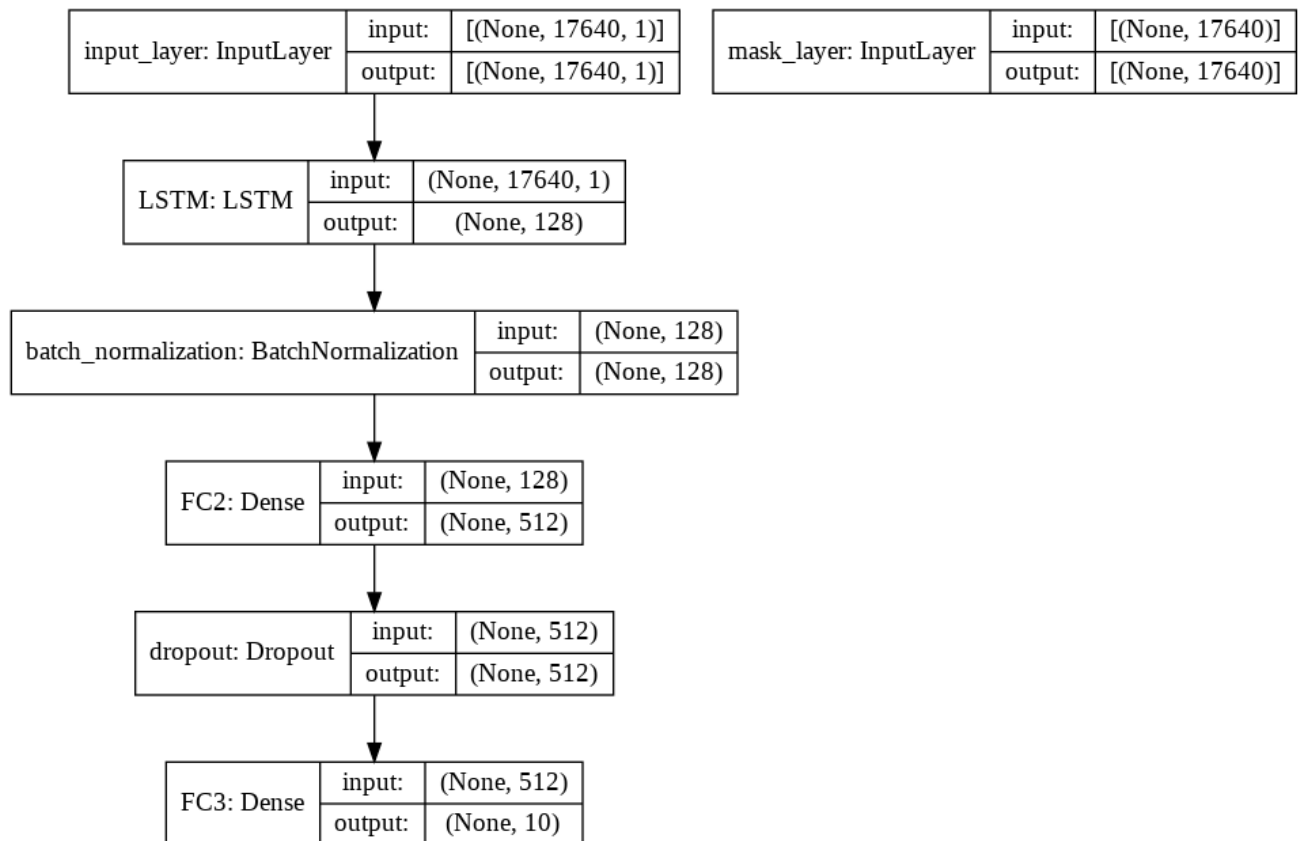
✕ .set\_ylabel('True labels')

```
Text(0.5, 1.0, 'Confusion Matrix')
```



```
plot_model(model_raw, show_shapes=True, show_layer_names=True)
```





Saved successfully!

## 2. (Model 2) Converting into spectrogram and giving spectrogram data as input.

We can use librosa to convert raw data into spectrogram. A spectrogram shows the features in a two-dimensional representation with the intensity of a frequency at a point in time i.e we are converting Time domain to frequency domain. You can read more about this in

<https://pnsn.org/spectrograms/what-is-a-spectrogram>

```
# Function to convert raw padded audio signal into spectrogram.
def convert_to_spectrogram(raw_data):
    '''converting to spectrogram'''
    spectrum = librosa.feature.melspectrogram(y=raw_data, sr=sample_rate, n_mels=64)
    logmel_spectrum = librosa.power_to_db(S=spectrum, ref=np.max)
    return logmel_spectrum
```

Use convert\_to\_spectrogram and convert every raw sequence in X\_train\_pad\_seq and X\_test\_pad\_seq.

Save those all in the X\_train\_spectrogram and X\_test\_spectrogram ( These two arrays must be numpy arrays)

```
X_train_spectrogram = []
for j in X_train_pad_seq:
    l = convert_to_spectrogram(j)
    X_train_spectrogram.append(l)
X_train_spectrogram = np.array(X_train_spectrogram)
```



```
X_test_spectrogram = []
for j in X_test_pad_seq:
    l = convert_to_spectrogram(j)
    X_test_spectrogram.append(l)
X_test_spectrogram = np.array(X_test_spectrogram)
```

```
plt.subplots_adjust(wspace=1, hspace=1)
for i in range(0,3):
    fig, ax = plt.subplots()
    img = librosa.display.specshow(convert_to_spectrogram(X_train_processed['raw_data'])[i])
    ax.set(title='Mel spectrogram display')
    ax.set_ylim([0,6000])
    fig.colorbar(img, ax=ax, format="%+2.f dB")
```

Saved successfully!




&lt;Figure size 432x288 with 0 Axes&gt;

**Grader function 6**

```
def grader_spectrogram():
    flag_shape = (X_train_spectrogram.shape==(1400,64, 35)) and (X_test_spectrogram.shape
    return flag_shape
grader_spectrogram()
```

True



Saved successfully!



X\_train

Test data: X\_test\_spectrogram and y\_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes "X\_train\_spectrogram" as input and has to return output at every time step.
2. Average the output of every time step and give this to the Dense layer of any size.
3. give the above output to Dense layer of size 10( output layer) and train the network with sparse categorical cross entropy.
4. Use tensorboard to plot the graphs of loss and metric(use micro F1 score as metric) and histograms of gradients.
5. make sure that it won't overfit.
6. You are free to include any regularization

```
tf.keras.backend.clear_session()
reg = tf.keras.regularizers.L2(l2=1.5)
```

```

tf.keras.backend.clear_session()
input_layer = Input(shape=(64,35), name = 'input_layer')
ls = LSTM(units = 128, name = 'LSTM', return_sequences = True)(input_layer)
ad = GlobalAveragePooling1D()(ls)
dc1 = Dense(1024,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed
dc1 = Dense(256,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed
dc1 = BatchNormalization()(dc1)
dc1 = Dropout(0.6899)(dc1)
out = Dense(10,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal

```

```
m_spec = Model(inputs = input_layer, outputs = out)
```

```
m_spec.summary()
```

Model: "model"

| Layer (type)                 | Output Shape     | Param # |
|------------------------------|------------------|---------|
| input_layer (InputLayer)     | [(None, 64, 35)] | 0       |
| LSTM (LSTM)                  | (None, 64, 128)  | 83968   |
| global_average_pooling1d (G1 | (None, 128)      | 0       |
| FC1 (Dense)                  | (None, 1024)     | 132096  |
| FC2 (Dense)                  | (None, 256)      | 262400  |
|                              | (None, 256)      | 1024    |
|                              | (None, 256)      | 0       |
| FC3 (Dense)                  | (None, 10)       | 2570    |
| Total params: 482,058        |                  |         |
| Trainable params: 481,546    |                  |         |
| Non-trainable params: 512    |                  |         |

Saved successfully!



```

save = 'model_spec_save/*.hdf5'
r = glob.glob(save)
for i in r:
    os.remove(i)

```

```

filepath="model_spec_save/model-{epoch:02d}-{val_sparse_categorical_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_sparse_categorical_accuracy',
reducelr = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.055, patience = 1, verbose =

```

```

lrschedule = LearningRateScheduler(changeLearningRate, verbose=1)
initial_learningrate=0.001

```

```
Train_data_spec = [X_train_spectrogram, y_train.values]
```

```

Test_data_spec = [X_test_spectrogram, y_test.values]

metrics = Metrics(Train_data_spec, Test_data_spec)

callbacks = [metrics, checkpoint, reduce_lr, lrschedule]

m_spec.compile(optimizer=tf.keras.optimizers.Adam(0.001),
               loss='sparse_categorical_crossentropy',
               metrics=['sparse_categorical_accuracy'])

tf.keras.backend.clear_session()
model_spec_history = m_spec.fit(X_train_spectrogram, y_train.values, batch_size=50, valid

```

Epoch 1/50

Epoch 00001: LearningRateScheduler reducing learning rate to 0.001.  
 28/28 [=====] - 2s 23ms/step - loss: 2.7975 - sparse\_categorical\_accuracy: 0.1583  
 - train\_f1\_score : 0.16857142857142857 - val\_f1\_score : 0.15833333333333333

Epoch 00001: val\_sparse\_categorical\_accuracy improved from -inf to 0.15833, saving model to model\_spec\_history.h5  
 Epoch 2/50

Epoch 00002: LearningRateScheduler reducing learning rate to 0.001.  
 28/28 [=====] - 0s 10ms/step - loss: 2.0661 - sparse\_categorical\_accuracy: 0.3250  
 - train\_f1\_score : 0.32571428571428573 - val\_f1\_score : 0.325

Epoch 00002: val\_sparse\_categorical\_accuracy improved from 0.15833 to 0.32500, saving model to model\_spec\_history.h5  
 Epoch 3/50

Saved successfully!

Epoch 00003: LearningRateScheduler reducing learning rate to 0.001.  
 28/28 [=====] - 0s 8ms/step - loss: 1.8349 - sparse\_categorical\_accuracy: 0.2128  
 - train\_f1\_score : 0.21285714285714286 - val\_f1\_score : 0.19

Epoch 00003: val\_sparse\_categorical\_accuracy did not improve from 0.32500

Epoch 00003: ReduceLROnPlateau reducing learning rate to 5.500000261235982e-05.  
 Epoch 4/50

Epoch 00004: LearningRateScheduler reducing learning rate to 0.001.  
 28/28 [=====] - 0s 8ms/step - loss: 1.5767 - sparse\_categorical\_accuracy: 0.2483  
 - train\_f1\_score : 0.28714285714285714 - val\_f1\_score : 0.24833333333333332

Epoch 00004: val\_sparse\_categorical\_accuracy did not improve from 0.32500  
 Epoch 5/50

Epoch 00005: LearningRateScheduler reducing learning rate to 0.00055.  
 28/28 [=====] - 0s 7ms/step - loss: 1.5362 - sparse\_categorical\_accuracy: 0.3066  
 - train\_f1\_score : 0.3507142857142857 - val\_f1\_score : 0.30666666666666664

Epoch 00005: val\_sparse\_categorical\_accuracy did not improve from 0.32500  
 Epoch 6/50

Epoch 00006: LearningRateScheduler reducing learning rate to 0.00055.  
 28/28 [=====] - 0s 7ms/step - loss: 1.3419 - sparse\_categorical\_accuracy: 0.3416  
 - train\_f1\_score : 0.3985714285714285 - val\_f1\_score : 0.34166666666666667

Epoch 00006: val\_sparse\_categorical\_accuracy improved from 0.32500 to 0.34167, saving Epoch 7/50

Epoch 00007: LearningRateScheduler reducing learning rate to 0.00055.  
28/28 [=====] - 0s 8ms/step - loss: 1.2128 - sparse\_categorical\_crossentropy : 0.5007142857142857 - val\_f1\_score : 0.42999999999999994

Epoch 00007: val\_sparse\_categorical\_accuracy improved from 0.34167 to 0.43000, saving Epoch 8/50

Epoch 00008: LearningRateScheduler reducing learning rate to 0.00055.  
28/28 [=====] - 0s 9ms/step - loss: 1.1470 - sparse\_categorical\_crossentropy : 0.495 - val\_f1\_score : 0.44166666666666665

```
acc = model_spec_history.history['sparse_categorical_accuracy']
val_acc = model_spec_history.history['val_sparse_categorical_accuracy']
```

```
loss = model_spec_history.history['loss']
val_loss = model_spec_history.history['val_loss']
```

```
f1 = metrics.history['train_f1_score']
val_f1 = metrics.history['val_f1_score']
```

```
epochs_range = range(50)
```

```
#plt.figure(figsize=(8, 8))
#plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
```


Saved successfully!

```
plt.show()
```


```
plt.plot(epochs_range, f1, label='Training micro F1 score')
plt.plot(epochs_range, val_f1, label='Validation micro F1 score')
plt.legend(loc='lower right')
plt.title('Training and Validation F1 Score')
plt.show()
```

```
#plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```


Training and Validation Accuracy



Training and Validation F1 Score



Training and Validation Loss



## as discussed above, please write the LSTM

```
y_pred = m_spec.predict(X_test_spectrogram)
```

```
y_pred = np.argmax(y_pred, axis = 1)
```

```
print('Classification Report')
```

```
print(classification_report(y_test.values, y_pred))
```

Classification Report


|   | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.88      | 0.95   | 0.91     | 60      |
| 1 | 0.85      | 0.87   | 0.86     | 60      |
| 2 | 0.84      | 0.82   | 0.83     | 60      |
| 3 | 0.68      | 0.67   | 0.67     | 60      |
| 4 | 0.87      | 0.90   | 0.89     | 60      |
| 5 | 0.89      | 0.78   | 0.83     | 60      |
| 6 | 0.70      | 0.78   | 0.74     | 60      |
| 7 | 0.70      | 0.78   | 0.74     | 60      |
| 8 | 0.81      | 0.72   | 0.76     | 60      |
| 9 | 0.85      | 0.78   | 0.82     | 60      |

|              |      |      |      |     |
|--------------|------|------|------|-----|
| accuracy     |      |      | 0.81 | 600 |
| macro avg    | 0.81 | 0.81 | 0.81 | 600 |
| weighted avg | 0.81 | 0.81 | 0.81 | 600 |

```
ax= plt.subplot()
sns.heatmap(confusion_matrix(y_test.values, y_pred), annot=True, ax = ax, fmt='g', cmap='G

# labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels')
ax.set_title('Confusion Matrix')
```

```
Text(0.5, 1.0, 'Confusion Matrix')
```

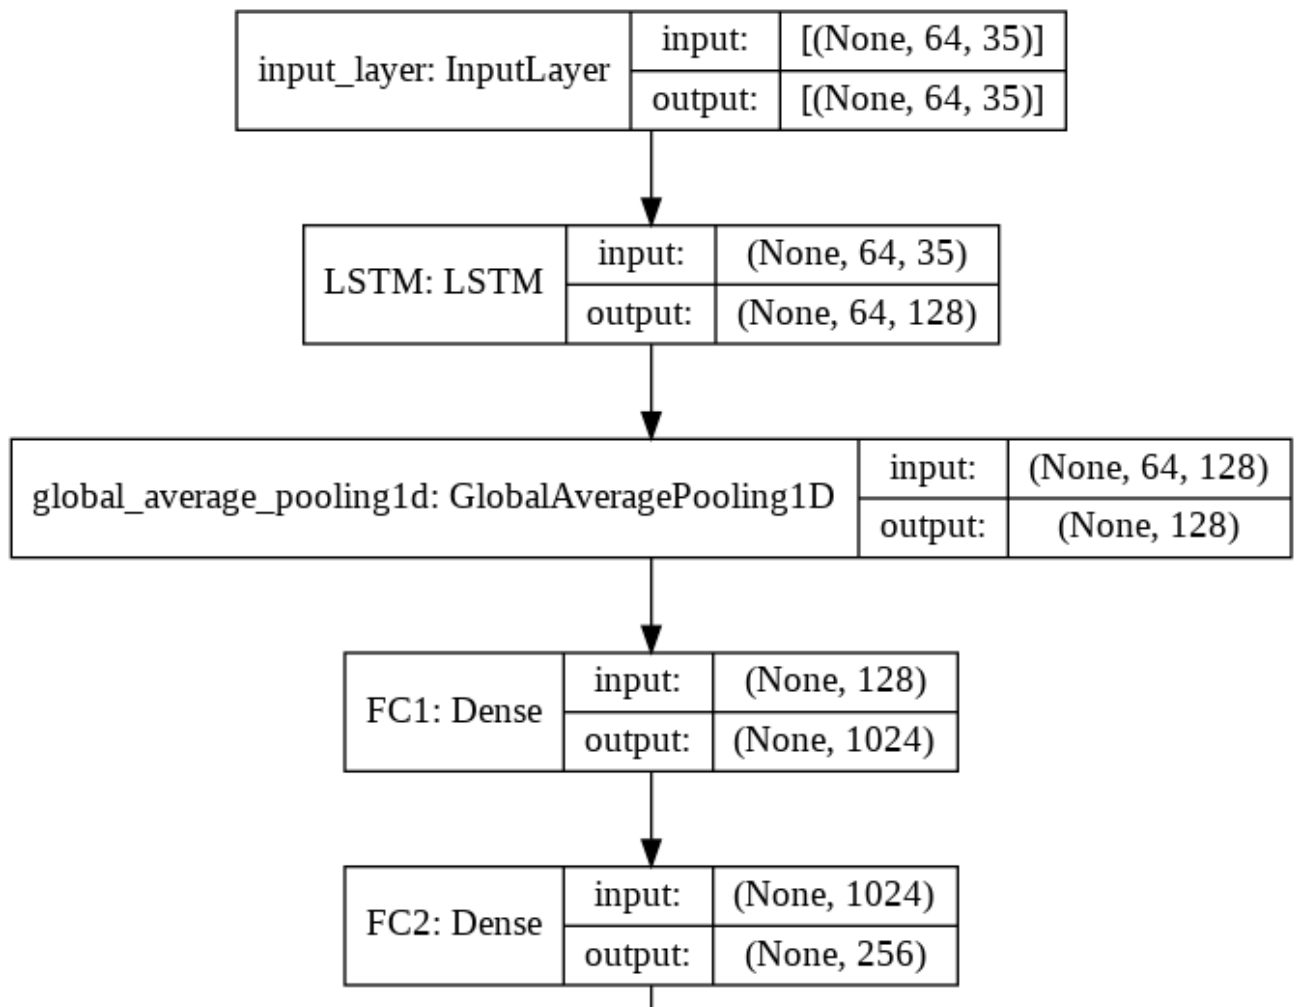


Saved successfully!

✕ al.h5')

```
m_spec = load_model('/content/Model_Spect_1_0.805_final.h5')
plot_model(m_spec, show_shapes=True, show_layer_names=True)
```





## ▼ Data Augmentation

Saved successfully!

| output: | (None, 256) |

I will now we have done with 2000 samples only. It is very less data. We are giving the process of generating augmented data below.

There are two types of augmentation:

1. time stretching - Time stretching either increases or decreases the length of the file. For time stretching we move the file 30% faster or slower
2. pitch shifting - pitch shifting moves the frequencies higher or lower. For pitch shifting we shift up or down one half-step.

```

## generating augmented data.
def generate_augmented_data(file_path, label):
    augmented_data = []
    samples = load_wav(file_path, get_duration=False)
    augmented_data.append(samples)
    for time_value in [0.7, 1, 1.3]:
        for pitch_value in [-1, 0, 1]:
            time_stretch_data = librosa.effects.time_stretch(samples, rate=time_value)
            final_data = librosa.effects.pitch_shift(time_stretch_data, sr=sample_rate, n_
            augmented_data.append(final_data)
    return augmented_data, np.full(len(augmented_data), label)

```

```
Y = df_audio['label'].astype("int32")
```

```
temp_path, temp_lab = df_audio.iloc[1].path, Y.iloc[1]
aug_temp1, auj= generate_augmented_data(temp_path, temp_lab)
len(aug_temp1)
```

```
10
```

As discussed above, for one data point, we will get 9 augmented data points and original point. We have 2000 data points(train plus test) so, after augmentation we will get 20000 ( train - 14000, test - 6000).

do the above steps i.e training with raw data and spectrogram data with augmentation.

```
a = Parallel(n_jobs=-1, verbose = 1)(delayed(generate_augmented_data)(i, j) for i, j in zi
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed:    7.9s
[Parallel(n_jobs=-1)]: Done 196 tasks    | elapsed:   42.3s
[Parallel(n_jobs=-1)]: Done 446 tasks    | elapsed:   1.8min
[Parallel(n_jobs=-1)]: Done 796 tasks    | elapsed:   3.2min
[Parallel(n_jobs=-1)]: Done 1246 tasks   | elapsed:   5.1min
[Parallel(n_jobs=-1)]: Done 1796 tasks   | elapsed:   7.4min
[Parallel(n_jobs=-1)]: Done 2000 out of 2000 | elapsed:   8.2min finished
```

Saved successfully!

```
New_labels = a[:,1].astype('int32').ravel()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWarning:
    """Entry point for launching an IPython kernel.
```

```
len(New_samples), len(New_labels)
```

```
(20000, 20000)
```

```
# Shuffling the above dataset.
```

```
New_samples, New_labels = shuffle(New_samples, New_labels, random_state=33)#don't change t
```

```
#split the data into train and validation and save in X_train, X_test, y_train, y_test
#use stratify sampling, use random state of 45, use test size of 30%
```

```
X_train, X_test, y_train, y_test = train_test_split(New_samples, New_labels, test_size = 0
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
((14000,), (6000,), (14000,), (6000,))
```

```
X_train_mask = np.array([np.ones(g.shape[0]) for g in X_train])
X_train_mask.shape
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWarning:
    """Entry point for launching an IPython kernel.
(14000,)
```

```
X_test_mask = np.array([np.ones(g.shape[0]) for g in X_test])
X_test_mask.shape
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: VisibleDeprecationWarning:
    """Entry point for launching an IPython kernel.
(6000,)
```

```
max_length = 17640
X_train_mask[-50].shape, X_train[-50].shape
```

```
((10232,), (10232,))
```

```
X_train_pad_seq = pad_sequences(X_train, maxlen=max_length, padding='post', dtype = np.float32)
X_train_mask = pad_sequences(X_train_mask, maxlen=max_length, padding='post', dtype = bool)
```

```
X_test_pad_seq = pad_sequences(X_test, maxlen=max_length, padding='post', dtype = np.float32)
X_test_mask = pad_sequences(X_test_mask, maxlen=max_length, padding='post', dtype = bool,
```

Saved successfully!

```
..., False, False, False],
[ True,  True,  True, ..., False, False, False],
[ True,  True,  True, ..., False, False, False],
...,
[ True,  True,  True, ..., False, False, False],
[ True,  True,  True, ..., False, False, False],
[ True,  True,  True, ..., False, False, False]]])
```

```
X_T_pad = tf.expand_dims(X_train_pad_seq, axis = 2)
X_T_pad.shape, X_T_pad[0].shape
```

```
(TensorShape([14000, 17640, 1]), TensorShape([17640, 1]))
```

```
X_Te_pad = tf.expand_dims(X_test_pad_seq, axis = 2)
X_Te_pad.shape, X_Te_pad[0].shape
```

```
(TensorShape([6000, 17640, 1]), TensorShape([17640, 1]))
```

```
td1 = tf.data.Dataset.from_tensor_slices((X_T_pad, X_train_mask))
td2 = tf.data.Dataset.from_tensor_slices((y_train))
td = tf.data.Dataset.zip((td1, td2))
```

```

BATCH_SIZE = 128
SHUFFLE_BUFFER_SIZE = 100
train_dataset = td.shuffle(SHUFFLE_BUFFER_SIZE).batch(BATCH_SIZE)

td1 = tf.data.Dataset.from_tensor_slices((X_Te_pad, X_test_mask))
td2 = tf.data.Dataset.from_tensor_slices((y_test))
td = tf.data.Dataset.zip((td1, td2))

test_dataset = td.shuffle(SHUFFLE_BUFFER_SIZE).batch(BATCH_SIZE)

Train = [X_T_pad, X_train_mask]
Val = [X_Te_pad, X_test_mask]

Train_data = (Train, y_train)

Val_data = (Val, y_test)

del a

```

## ▼ 1. (Model 3) Giving Raw data directly.

```

tf.keras.backend.clear_session()
class Metrics(tf.keras.callbacks.Callback):
    def on_train_begin(self, logs = {}):
        self.training_data = x
        self.validation_data = y

    def on_train_end(self, logs = {}):
        None):

    def on_train_begin(self, logs = {}):
        ## on begin of training, we are creating a instance variable called history
        self.history={'train_f1_score': [], 'val_f1_score': []}

    def on_epoch_end(self, epoch, logs = {}):
        ## on end of each epoch, we will get logs and update the self.history dict
        train_predict = self.model.predict(self.training_data[0], batch_size = 128)
        train_bin = np.argmax(train_predict, axis = 1)
        train_targ = self.training_data[1]
        _train_f1 = f1_score(train_targ, train_bin, average = 'micro')

        val_predict = self.model.predict(self.validation_data[0], batch_size = 128)
        val_bin = np.argmax(val_predict, axis = 1)
        val_targ = self.validation_data[1]
        _val_f1 = f1_score(val_targ, val_bin, average = 'micro')

        self.history['val_f1_score'].append(_val_f1)
        self.history['train_f1_score'].append(_train_f1)

```

```
print(' - train_f1_score : ', _train_f1, ' - val_f1_score : ', _val_f1)
return
```

```
def changeLearningRate(epoch):
    global initial_learningrate
    epoch = epoch + 1
    if epoch % 5 == 0:
        initial_learningrate *= 0.55
    return initial_learningrate
```

```
import os
save = 'model_aug_raw_save/*.hdf5'
r = glob.glob(save)
for i in r:
    os.remove(i)
```

```
filepath="model_aug_raw_save/model-{epoch:02d}-{val_sparse_categorical_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_sparse_categorical_accuracy',
reducelr = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.001, patience = 1, verbose =
```

```
lrschedule = LearningRateScheduler(changeLearningRate, verbose=1)
initial_learningrate=0.001
```

```
metrics = Metrics(Train_data, Val_data)
```

Saved successfully!



```
reducelr, lrschedule]
```

```
reg = tf.keras.regularizers.L2(l2=0.01)
```

```
tf.keras.backend.clear_session()
input_layer = Input(shape=(17640,1), name = 'input_layer')
input_mask = Input(shape=(17640,), name = 'mask_layer', dtype=bool)
ls = LSTM(units = 128, name = 'LSTM')(input_layer, mask = input_mask)
ls = BatchNormalization()(ls)
dc1 = Dense(512,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed
dc1 = Dropout(0.35)(dc1)
out = Dense(10,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal
```

```
model_aug_raw = Model(inputs = [input_layer,input_mask], outputs = out)
```

```
model_aug_raw.summary()
```

Model: "model"

| Layer (type) | Output Shape | Param # | Connected to |
|--------------|--------------|---------|--------------|
| =====        |              |         |              |

|                                     |                    |       |                                       |
|-------------------------------------|--------------------|-------|---------------------------------------|
| input_layer (InputLayer)            | [(None, 17640, 1)] | 0     |                                       |
| mask_layer (InputLayer)             | [(None, 17640)]    | 0     |                                       |
| LSTM (LSTM)                         | (None, 128)        | 66560 | input_layer[0][0]<br>mask_layer[0][0] |
| batch_normalization (BatchNormaliza | (None, 128)        | 512   | LSTM[0][0]                            |
| FC2 (Dense)                         | (None, 512)        | 66048 | batch_normalization[                  |
| dropout (Dropout)                   | (None, 512)        | 0     | FC2[0][0]                             |
| FC3 (Dense)                         | (None, 10)         | 5130  | dropout[0][0]                         |
| =====                               |                    |       |                                       |
| Total params: 138,250               |                    |       |                                       |
| Trainable params: 137,994           |                    |       |                                       |
| Non-trainable params: 256           |                    |       |                                       |

```
model_aug_raw.compile(optimizer=tf.keras.optimizers.RMSprop(0.001),
                      loss='sparse_categorical_crossentropy',
                      metrics=['sparse_categorical_accuracy'])
```

```
train_steps = X_train.shape[0]//100
valid_steps = X_test.shape[0]//100
train_steps
```

140

Saved successfully!



```
train_dataset, validation_data = test_dataset, epochs =
```

Epoch 00007: val\_sparse\_categorical\_accuracy did not improve from 0.11983

Epoch 00007: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.  
Epoch 8/50

Epoch 00008: LearningRateScheduler reducing learning rate to 0.00055.  
110/110 [=====] - 112s 1s/step - loss: 3.1387 - sparse\_cat  
- train\_f1\_score : 0.10064285714285715 - val\_f1\_score : 0.09966666666666667

Epoch 00008: val\_sparse\_categorical\_accuracy did not improve from 0.11983

Epoch 00008: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.  
Epoch 9/50

Epoch 00009: LearningRateScheduler reducing learning rate to 0.00055.  
110/110 [=====] - 112s 1s/step - loss: 2.2875 - sparse\_cat  
- train\_f1\_score : 0.103 - val\_f1\_score : 0.10233333333333333

Epoch 00009: val\_sparse\_categorical\_accuracy did not improve from 0.11983

Epoch 00009: ReduceLROnPlateau reducing learning rate to 5.499999970197678e-07.  
Epoch 10/50

Epoch 00010: LearningRateScheduler reducing learning rate to 0.0003025000000000000

```
110/110 [=====] - 112s 1s/step - loss: 2.2450 - sparse_categorical_accuracy: 0.1055
- train_f1_score : 0.10214285714285715 - val_f1_score : 0.1055
```

Epoch 00010: val\_sparse\_categorical\_accuracy did not improve from 0.11983

Epoch 00010: ReduceLROnPlateau reducing learning rate to 3.0250000418163836e-07.  
Epoch 11/50

```
Epoch 00011: LearningRateScheduler reducing learning rate to 0.000302500000000000
110/110 [=====] - 111s 1s/step - loss: 2.2433 - sparse_categorical_accuracy: 0.1055
- train_f1_score : 0.10885714285714285 - val_f1_score : 0.10833333333333334
```

Epoch 00011: val\_sparse\_categorical\_accuracy did not improve from 0.11983

Epoch 00011: ReduceLROnPlateau reducing learning rate to 3.0250000418163836e-07.  
Epoch 12/50

```
Epoch 00012: LearningRateScheduler reducing learning rate to 0.000302500000000000
110/110 [=====] - 112s 1s/step - loss: 2.2810 - sparse_categorical_accuracy: 0.117
- train_f1_score : 0.11271428571428571 - val_f1_score : 0.117
```

Epoch 00012: val\_sparse\_categorical\_accuracy did not improve from 0.11983

Epoch 00012: ReduceLROnPlateau reducing learning rate to 3.0250000418163836e-07.  
Epoch 13/50

```
Epoch 00013: LearningRateScheduler reducing learning rate to 0.000302500000000000
110/110 [=====] - 112s 1s/step - loss: 2.2533 - sparse_categorical_accuracy: 0.1173
- train_f1_score : 0.121 - val_f1_score : 0.11733333333333333
```

Epoch 00013: val\_sparse\_categorical\_accuracy did not improve from 0.11983

Epoch 00013: ReduceLROnPlateau reducing learning rate to 3.0250000418163836e-07.

Saved successfully!

```
acc = model_history.history['sparse_categorical_accuracy']
val_acc = model_history.history['val_sparse_categorical_accuracy']
```

```
loss = model_history.history['loss']
val_loss = model_history.history['val_loss']
```

```
f1 = metrics.history['train_f1_score']
val_f1 = metrics.history['val_f1_score']
```



```
epochs_range = range(50)
```

```
#plt.figure(figsize=(8, 8))
#plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.show()
```

```
plt.plot(epochs_range, f1, label='Training micro F1 score')
plt.plot(epochs_range, val_f1, label='Validation micro F1 score')
plt.legend(loc='lower right')
```

```
plt.title('Training and Validation F1 Score')
plt.show()

#plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



Saved successfully!



```
model_aug_raw.save('Model_Aug_raw_2_0.1795_final.h5')
```



```
model_raw = load_model('/content/model-31-0.1333.hdf5')
```

```
y_pred = model_aug_raw.predict(Val)
y_pred = np.argmax(y_pred, axis = 1)
```

```
print('Classification Report')
print(classification_report(y_test, y_pred))
```

```

Classification Report
              precision    recall  f1-score   support

     0       0.21       0.14       0.17        600
     1       0.10       0.01       0.03        600
     2       0.16       0.12       0.14        600
     3       0.24       0.23       0.24        600
     4       0.15       0.12       0.13        600
     5       0.13       0.06       0.08        600
     6       0.11       0.09       0.10        600
     7       0.17       0.11       0.13        600
     8       0.19       0.83       0.30        600
     9       0.23       0.09       0.13        600

 accuracy          0.18        6000
 macro avg         0.17        6000
 weighted avg      0.17        6000

```


```
ax= plt.subplot()
```

Saved successfully!

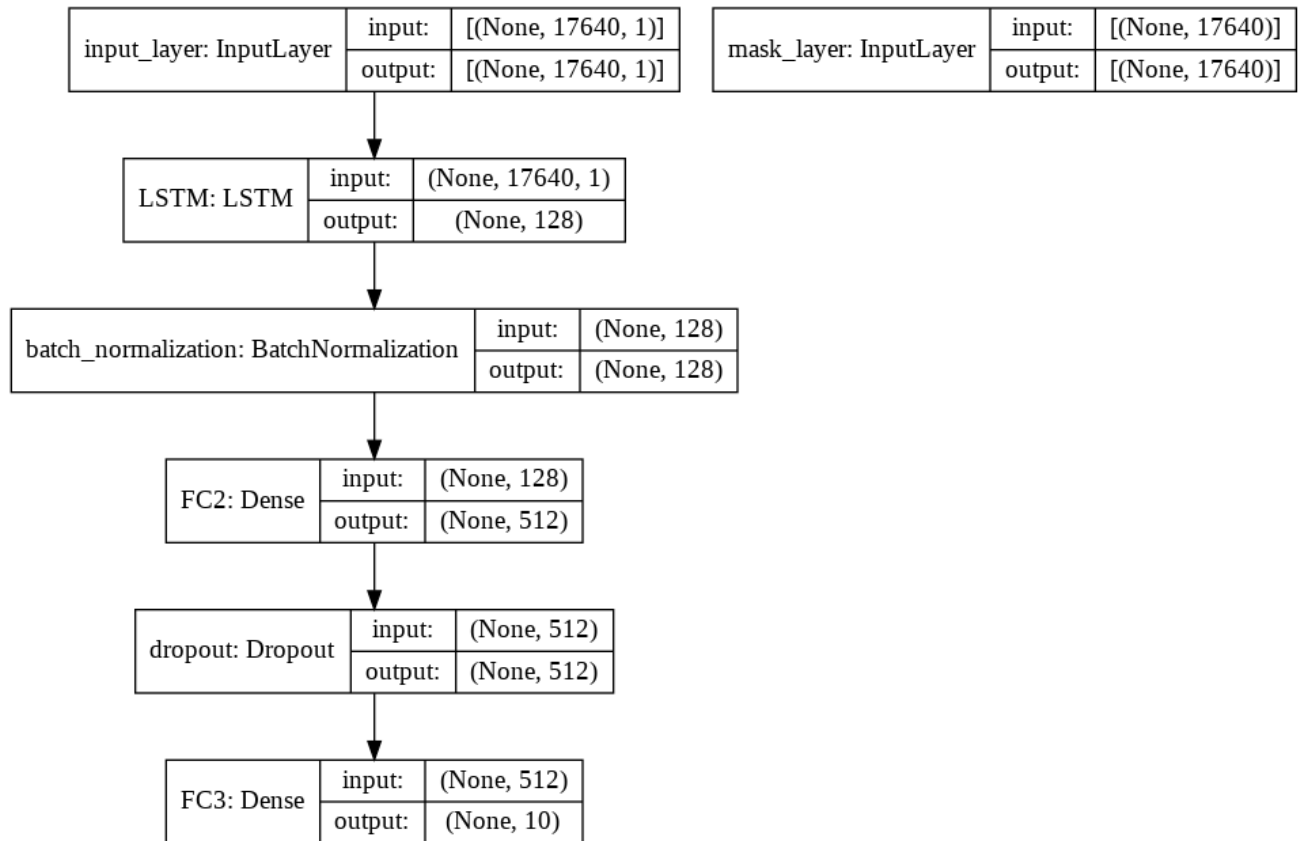
```
y_pred), annot=True, ax = ax, fmt='g', cmap='Greens')
```

```
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels')
ax.set_title('Confusion Matrix')
```

```
Text(0.5, 1.0, 'Confusion Matrix')
```



```
from tensorflow.keras.utils import plot_model
plot_model(model_aug_raw, show_shapes=True, show_layer_names=True)
```



Saved successfully!



## 2. (Model 4) Converting into spectrogram and giving spectrogram data as input.

We can use librosa to convert raw data into spectrogram. A spectrogram shows the features in a two-dimensional representation with the intensity of a frequency at a point in time i.e we are converting Time domain to frequency domain. you can read more about this in

<https://pnsn.org/spectrograms/what-is-a-spectrogram>

```
def convert_to_spectrogram(raw_data):
    '''converting to spectrogram'''
    spectrum = librosa.feature.melspectrogram(y=raw_data, sr=sample_rate, n_mels=64)
```

```
logmel_spectrum = librosa.power_to_db(S=spectrum, ref=np.max)
```

Use `convert_to_spectrogram` and convert every raw sequence in `X_train_pad_seq` and `X_test_pad_seq`.

Save those all in the `X_train_spectrogram` and `X_test_spectrogram` ( These two arrays must be numpy arrays)

```
X_train_spectrogram = []
for j in X_train_pad_seq:
    l = convert_to_spectrogram(j)
    X_train_spectrogram.append(l)
X_train_spectrogram = np.array(X_train_spectrogram)
```

```
X_test_spectrogram = []
for j in X_test_pad_seq:
    l = convert_to_spectrogram(j)
    X_test_spectrogram.append(l)
X_test_spectrogram = np.array(X_test_spectrogram)
```

## Grader function 7

```
def grader_spectrogram():
    flag_shape = (X_train_spectrogram.shape==(14000,64, 35)) and (X_test_spectrogram.shape
    return flag_shape
```

Saved successfully!



```
td1 = tf.data.Dataset.from_tensor_slices((X_train_spectrogram, y_train))
```

```
BATCH_SIZE = 128
```

```
SHUFFLE_BUFFER_SIZE = 100
```

```
train_spec_dataset = td1.shuffle(SHUFFLE_BUFFER_SIZE).batch(BATCH_SIZE)
```

```
td1 = tf.data.Dataset.from_tensor_slices((X_test_spectrogram, y_test))
```

```
test_spec_dataset = td1.shuffle(SHUFFLE_BUFFER_SIZE).batch(BATCH_SIZE)
```

Now we have

Train data: `X_train_spectrogram` and `y_train`

Test data: `X_test_spectrogram` and `y_test`

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes "X\_train\_spectrogram" as input and has to return output at every time step.
2. Average the output of every time step and give this to the Dense layer of any size.
3. give the above output to Dense layer of size 10( output layer) and train the network with sparse categorical cross entropy.
4. Use tensorboard to plot the graphs of loss and metric(use micro F1 score as metric) and histograms of gradients.
5. make sure that it won't overfit.
6. You are free to include any regularization

```
reg = tf.keras.regularizers.L2(l2=1.5)
```

```
tf.keras.backend.clear_session()
input_layer = Input(shape=(64,35), name = 'input_layer')
ls = LSTM(units = 128, name = 'LSTM', return_sequences = True)(input_layer)
ad = GlobalAveragePooling1D()(ls)
dc1 = Dense(1024,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed
dc1 = Dense(256,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed
dc1 = BatchNormalization()(dc1)
dc1 = Dropout(0.6599)(dc1)
out = Dense(10,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal
```

Saved successfully!

layer, outputs = out)

```
m_aug_spec.summary()
```

Model: "model"

| Layer (type)                 | Output Shape     | Param # |
|------------------------------|------------------|---------|
| input_layer (InputLayer)     | [(None, 64, 35)] | 0       |
| LSTM (LSTM)                  | (None, 64, 128)  | 83968   |
| global_average_pooling1d (Gl | (None, 128)      | 0       |
| FC1 (Dense)                  | (None, 1024)     | 132096  |
| FC2 (Dense)                  | (None, 256)      | 262400  |
| batch_normalization (BatchNo | (None, 256)      | 1024    |
| dropout (Dropout)            | (None, 256)      | 0       |
| FC3 (Dense)                  | (None, 10)       | 2570    |
| Total params: 482,058        |                  |         |

Trainable params: 481,546  
Non-trainable params: 512

---

```
save = 'model_aug_spec_save/*.hdf5'
r = glob.glob(save)
for i in r:
    os.remove(i)
```

```
filepath="model_aug_spec_save/model-{epoch:02d}-{val_sparse_categorical_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_sparse_categorical_accuracy',
reducelr = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.055, patience = 1, verbose =
```

```
lrschedule = LearningRateScheduler(changeLearningRate, verbose=1)
initial_learningrate=0.001
```

```
Train_data_spec = [X_train_spectrogram, y_train]
Test_data_spec = [X_test_spectrogram, y_test]
```

```
metrics = Metrics(Train_data_spec, Test_data_spec)
```

```
callbacks = [metrics, checkpoint, reducelr, lrschedule]
```

```
m_aug_spec.compile(optimizer=tf.keras.optimizers.Adam(0.001),
                    loss='sparse_categorical_crossentropy',
                    metrics=['sparse_categorical_accuracy'])
```

Saved successfully!



```
train_spec_dataset, validation_data = test_spec_dataset
```

```
Epoch 00006: LearningRateScheduler reducing learning rate to 0.00055.
110/110 [=====] - 1s 10ms/step - loss: 0.8274 - sparse_categorical_accuracy: 0.7166666666666667
- train_f1_score : 0.7325 - val_f1_score : 0.7166666666666667
```

```
Epoch 00006: val_sparse_categorical_accuracy did not improve from 0.72750
```

```
Epoch 00006: ReduceLROnPlateau reducing learning rate to 3.0249999836087226e-05.
Epoch 7/50
```

```
Epoch 00007: LearningRateScheduler reducing learning rate to 0.00055.
110/110 [=====] - 1s 10ms/step - loss: 0.7762 - sparse_categorical_accuracy: 0.7030714285714286
- train_f1_score : 0.7030714285714286 - val_f1_score : 0.688
```

```
Epoch 00007: val_sparse_categorical_accuracy did not improve from 0.72750
```

```
Epoch 00007: ReduceLROnPlateau reducing learning rate to 3.0249999836087226e-05.
Epoch 8/50
```

```
Epoch 00008: LearningRateScheduler reducing learning rate to 0.00055.
110/110 [=====] - 1s 10ms/step - loss: 0.7559 - sparse_categorical_accuracy: 0.7588333333333333
- train_f1_score : 0.7702857142857142 - val_f1_score : 0.7588333333333333
```

```
Epoch 00008: val_sparse_categorical_accuracy improved from 0.72750 to 0.75883, saving model to model_aug_spec_save/model-02-0.75883.hdf5
Epoch 9/50
```

```

Epoch 00009: LearningRateScheduler reducing learning rate to 0.00055.
110/110 [=====] - 1s 10ms/step - loss: 0.7339 - sparse_categorical_accuracy: 0.7583
- train_f1_score : 0.7 - val_f1_score : 0.692

Epoch 00009: val_sparse_categorical_accuracy did not improve from 0.75883

Epoch 00009: ReduceLROnPlateau reducing learning rate to 3.0249999836087226e-05.
Epoch 10/50

Epoch 00010: LearningRateScheduler reducing learning rate to 0.00030250000000000000
110/110 [=====] - 1s 10ms/step - loss: 0.6635 - sparse_categorical_accuracy: 0.7888
- train_f1_score : 0.8067142857142857 - val_f1_score : 0.7888333333333334

Epoch 00010: val_sparse_categorical_accuracy improved from 0.75883 to 0.78883, saving model to model_00010.h5
Epoch 11/50

Epoch 00011: LearningRateScheduler reducing learning rate to 0.00030250000000000000
110/110 [=====] - 1s 10ms/step - loss: 0.6383 - sparse_categorical_accuracy: 0.7928
- train_f1_score : 0.8132857142857143 - val_f1_score : 0.7928333333333333

Epoch 00011: val_sparse_categorical_accuracy improved from 0.78883 to 0.79283, saving model to model_00011.h5
Epoch 12/50

Epoch 00012: LearningRateScheduler reducing learning rate to 0.00030250000000000000
110/110 [=====] - 1s 10ms/step - loss: 0.6282 - sparse_categorical_accuracy: 0.8018
- train_f1_score : 0.8192142857142857 - val_f1_score : 0.8018333333333333

Epoch 00012: val_sparse_categorical_accuracy improved from 0.79283 to 0.80183, saving model to model_00012.h5
Epoch 13/50

Epoch 00013: LearningRateScheduler reducing learning rate to 0.00030250000000000000
110/110 [=====] - 1s 10ms/step - loss: 0.5982 - sparse_categorical_accuracy: 0.8118
- train_f1_score : 0.8257142857142857 - val_f1_score : 0.8118333333333333

Epoch 00013: val_sparse_categorical_accuracy improved from 0.80183 to 0.81183, saving model to model_00013.h5
Epoch 14/50

acc = model_spec_history.history['sparse_categorical_accuracy']
val_acc = model_spec_history.history['val_sparse_categorical_accuracy']

loss = model_spec_history.history['loss']
val_loss = model_spec_history.history['val_loss']

f1 = metrics.history['train_f1_score']
val_f1 = metrics.history['val_f1_score']

epochs_range = range(50)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.show()

plt.plot(epochs_range, f1, label='Training micro F1 score')
plt.plot(epochs_range, val_f1, label='Validation micro F1 score')
plt.legend(loc='lower right')
plt.title('Training and Validation F1 Score')



```

Saved successfully!




```
plt.show()
```

```
#plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



Saved successfully!



```
## as discussed above, please write the LSTM
y_pred = m_aug_spec.predict(X_test_spectrogram)
```

```

y_pred = np.argmax(y_pred, axis = 1)
#p = p.round()

print('Classification Report')
print(classification_report(y_test, y_pred))

```

```

Classification Report
              precision    recall  f1-score   support

     0       0.95         0.94         0.95         600
     1       0.91         0.88         0.90         600
     2       0.82         0.84         0.83         600
     3       0.75         0.71         0.73         600
     4       0.96         0.98         0.97         600
     5       0.90         0.89         0.89         600
     6       0.80         0.78         0.79         600
     7       0.83         0.86         0.84         600
     8       0.83         0.85         0.84         600
     9       0.90         0.92         0.91         600

 accuracy          0.87          0.87          0.87        6000
  macro avg       0.87          0.87          0.87        6000
 weighted avg     0.87          0.87          0.87        6000

```


```

ax= plt.subplot()
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, ax = ax, fmt='g', cmap='Greens')

# labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels')

```

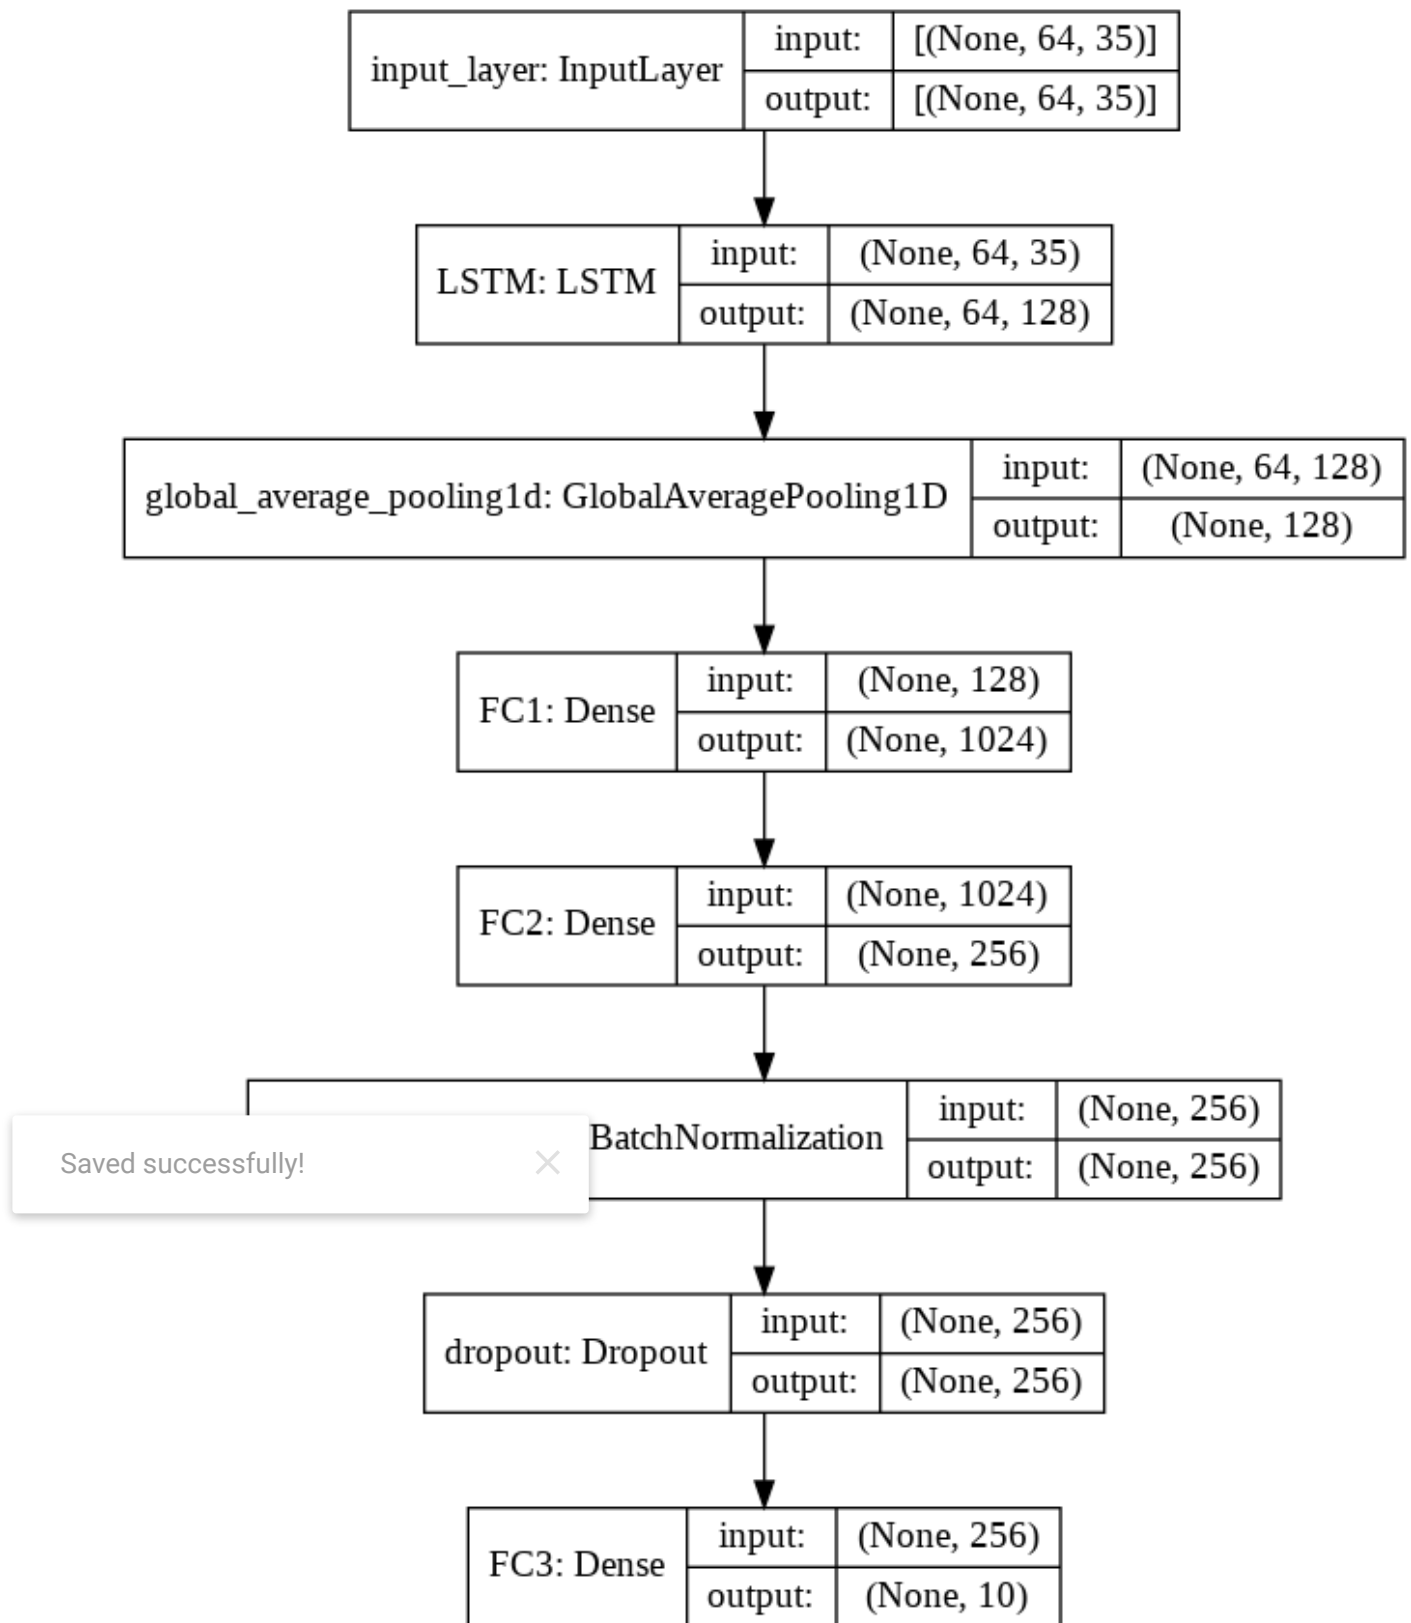
Saved successfully!



```
m_aug_spec.save('Model_Aug_Spect_2_0.805_final.h5')
```

```
plot_model(m_aug_spec, show_shapes=True, show_layer_names=True)
```





## Observations:

| Data                    | Model     | Test_F1_Score |
|-------------------------|-----------|---------------|
| Original - Raw Data     | Model - 1 | 0.1517        |
| Original - Spectrogram  | Model - 2 | 0.805         |
| Augmented - Raw Data    | Model - 3 | 0.1795        |
| Augmented - Spectrogram | Model - 4 | 0.8658        |

1. Micro F1 score for model 1 is: 0.15175246013
2. Micro F1 score for model 2 is: 0.80541004575
3. Micro F1 score for model 2 is: 0.17958463700
4. Micro F1 score for model 2 is: 0.86580286491

Spectrogram data gives good F1 score than using raw data.

Saved successfully!



✓ 0s completed at 12:33 AM

